

Some Reminders for a Seamless Online Class...

- Please turn on your video
- Mute yourself (press and hold spacebar when you'd like to talk)
- Don't do anything you wouldn't do in an in-person class
- I will occasionally check the chat for messages if you'd like to share there instead
- Please say your name before you speak



Recap

- Data-savviness is the future!
- “Classical” relational databases
 - Notion of a DBMS
 - The relational data model and algebra: bags and sets
 - SQL Queries, Modifications, DDL
 - Database Design
 - Views, constraints, triggers, and indexes
 - Query processing & optimization
 - Transactions
- Non-classical data systems
 - Data preparation:
 - Semi-structured data and document stores
 - Unstructured data and search engines
 - Data Exploration:
 - Cell-structured data and spreadsheets
 - Dataframes and dataframe systems
 - OLAP, summarization, and visual analytics
 - Batch Analytics:
 - Compression and column stores
 - **Parallel data processing and map-reduce**



Parallel data processing and map-reduce

- We've studied OLAP — a specialization of relational databases targeted at business analytics and reporting at scale with data cube materialization and column stores
- Today, we're going to be studying the primitives for processing large volumes of relational, unstructured, or semi-structured data at scale
- Often, when we're trying to process really large volumes of data, we need to span across multiple nodes/machines
 - This hasn't been a focus of our class so far
- We'll start by covering what map-reduce offers, before switching over to cover parallel databases



Let's revisit search engines

- To create an inverted index, we need to read all the webpages
 - Size of the web: 20+B web pages x 20 KB = 400+ TB
 - Disk reading speed ~50MBPS
 - On a single disk drive: $(400 \times 1000 \times 1000) / (35 \times 60 \times 60 \times 24 \times 30)$
 - 4 months to read the web!
 - If each hard-drive can store 1TB, then 400 hard-drives to store the web.
- Instead, parallelize!
 - Modern data system architectures use many cheap (commodity) machines connected by cheap network (ethernet)
- Q: If we could read from all 400 hard-drives at once, we could get the job done in how much time?
 - $120 \times 24 \text{ hard-drive-hours} / 400 \text{ hard-drives} = 7.2 \text{ hours!}$




Large Scale Data Processing

- Challenges:
 - How do we specify the processing task?
 - How do we distribute this processing across these many machines?
 - Transferring data across network is expensive
 - How do we deal with failures?
 - If a server dies once in a 100 days, then if we have a 1000 servers, on average 10 will die per day



Google's Approach: Circa '00s

- First component: a distributed, replicated file system
 - GFS, modern open-source incarnation: HDFS
 - Just like a file system in your machine, but each file is replicated across machines for reliability
 - Files are very large; rarely updated, but reading and appending are more common
 - There is a “coordinator” node that keeps track of which file is where



- Second component: a distributed data processing programming paradigm
 - Map-Reduce, modern open-source incarnation
 - The processing usually happens “close” to the data if possible to avoid moving data across network (but in some cases unavoidable)



Map-Reduce

- Let's say we have the entire web crawl and we want to count the # of times each word appears on the web
 - The web is stored across 400 machines, each with its hard-drive
- Q: How would you go about doing this?



Map-Reduce Overview


- Developer-provided Processing Step 1: **Map**
 - Read a lot of data and extract something of value
- Internal Step 2: Group by Keys (**Shuffle**)
- Developer-provided Processing Step 2: **Reduce**
 - Aggregate, summarize, filter, transform data organized by key



Map-Reduce Overview

- Developer-provided Processing Step 1: **Map**
 - Read a lot of data and extract something of value
 - *For each document, emit $\langle \text{word}: 1 \rangle$ pairs — key: value pairs*
- Internal Step 2: Group by Keys (**Shuffle**)
 - *Group/sort pairs based on word, so: $\langle \text{word } 1, 1 \rangle, \langle \text{word } 1, 1 \rangle, \dots \langle \text{word } 1, 1 \rangle$
 $\langle \text{word2}, 1 \rangle \dots$*
- Developer-provided Processing Step 2: **Reduce**
 - Aggregate, summarize, filter, transform data organized by key
 - *For each word, sum up the total number of 1s in the value*





The Map-Reduce Paradigm is General!

- Simply change the map and reduce functions to fit new applications
- Semantics:
 - Programmer specifies two methods:
 - $Map(k, v) \rightarrow \langle k', v' \rangle^*$
 - Takes a key-value pair and outputs a set of key-value pairs
 - E.g., key is the filename, value is a single line in the file
 - There is one Map call for every (k,v) pair
 - $Reduce(k', \langle v' \rangle^*) \rightarrow \langle k', v'' \rangle^*$
 - All values v' with same key k' are reduced together and processed in v' order
 - There is one Reduce function call per unique key k'




For Our Word Count Example


- Map (key, value)
 - //key is documentname, value is contents
 - For each word in value
 - Output (word, 1)
- Reduce (key, values)
 - // key is word, values is a list of 1s
 - Output (key, size (values))



What conceptually happens



What actually happens



What does the Map-Reduce Paradigm Handle?

- Assigning tasks to machines
- Performing the “shuffle” step
- Handling failures and inter-machine communication
 - Does so via materialization after every step
- In some cases, can apply a “combiner” at the map task itself to reduce the amount of data shuffled, especially if the reduce function has some nice commutative properties
 - For example, we can sum up the 1s for each word for each document before shuffling across network



A slightly harder case: inverted index construction... remember this slide?

- Phase 1:
 - For each document, generate a “canonicalized term, documentID” pair for all non-stop-word terms in the document
 - Can encode more information with the pair, e.g., list of locations, frequency of occurrence, etc.
- Phase 2:
 - Sort these pairs based on the term, documentID pair
 - May need to use two or multi-pass sort (remember query processing!)
- Phase 3:
 - Read in the sorted pairs based on term, concatenate documentID together to form a posting list for that term
 - (We may be able to merge Phase 3 with Phase 2)



Q: How would we use M-R to construct the inverted index?

- Obama -> 10, 24, 125, 259, 1025, 2314, ...
- Bush -> 10, 15, 17, 259, 2001, 2547, ...
- Clinton -> 10, 15, 24, 17, 125, 1005, 2001, 2347, ...



Answer

- Map
 - Input: document
 - Output: <word: documentID>
 - Optionally, position
- Shuffle based on word
- Reduce
 - Input: <word, documentID list>
 - Output: <word, sorted list of documentIDs, dropping duplicates>
- What would a combiner do here?



Other Map-Reduce Exercises

- Compute the total degree (indegree plus outdegree) for each node in a graph
- Input directed graph representation:
 - “Edge-pair” representation
 - $G(V1, V2)$ split across many machines
- Q: how would we do this?
- Map:
 - Input $\langle V1, V2 \rangle$
 - Output $\langle V1: 1 \rangle, \langle V2: 1 \rangle$
- Reduce:
 - Input $\langle V1: \text{list of } 1s \rangle$
 - Output: $\langle V1: \text{size of list} \rangle$



Other Map-Reduce Exercises (II)

- Computing a natural join
- Input relations $R(A, B)$, $S(B, C)$ partitioned across many nodes
- Q: how would we do this?
- Map: for input R (S)
 - Output $\langle B: (R, A) \rangle$ ($\langle B: (S, C) \rangle$)
- Reduce: for input $\langle B: (R, A_1) \dots (R, A_n), (S, C_1) \dots (S, C_m) \rangle$
 - Output: cross product of A_i and C_j
 - $(A_1, B, C_1), (A_1, B, C_2), \dots, (A_n, B, C_m)$



Try at home!

- Q: How would you compute connected components in a graph?
- Q: How would you compute page rank?



How would we use M-R?

- Since Google introduced its Map-Reduce in the early 2000s, open source alternatives have appeared:
 - Hadoop MapReduce (2006) — still exists to this day for open-ended large scale data processing
- Other document stores also support M-R on json-type documents
 - Key-value stores are a natural fit for M-R
 - CouchDB
 - MongoDB (demo!)



Downsides of Map-Reduce?

- Q: what are potential downsides of map-reduce?
- Hand-writing algorithms: no indexing, no query optimization
- Materialization after every step (no pipelining)
- No “declarative” query processing



From “A Comparison of Approaches to Large-Scale Data Analysis”, SIGMOD’09

- Compares Vertica, an industry relational DBMS (DBMS-X), and Hadoop
- A few charts...
 - Since Hadoop writes out partial results locally, the “white” stacked bar shows the time to assemble the final result



Loading and Filtering




Figure 3: Load Times – UserVisits Data Set (20GB/node)




Figure 6: Selection Task Results



Aggregation and Joins




Figure 8: Aggregation Task Results (2,000 Groups)




Figure 9: Join Task Results



Moral of the Story

- While convenient to do arbitrary large scale parallel data processing on arbitrary data, Map-Reduce isn't actually all that efficient compared to parallel implementations of relational databases (column or row stores)
 - So if your data is relational, relational databases is definitely a better choice.
- So how do column and row stores work in parallel anyway?



Parallelism

- We've seen pipelined parallelism — where different operators work “in parallel” with each other, pipelining results
- Also can do “partitioned parallelism”
 - Where each operator has many parallel threads or processes reading data in parallel across partitions
 - These partitions could be in the same machine or across multiple machines



How does one partition the data?

- Vertical partitioning
 - Partitioning by columns
 - Very natural in column stores, certain column groups can be stored
- Horizontal partitioning
 - Partitioning by rows
 - e.g., rows A-M by name are in node 1, N-Z are in node 2
 - Various ways of doing this with trade-offs
 - Hash-based partitioning, e.g., based on $h(A)$
 - Range-based partitioning, e.g., A from [1, 100] in node 1, [101, 200] in node 2
 - Round-robin partitioning, add a tuple to each partition in sequence
 - Q: pros/cons?
 - Hashing is good to allow for all tuples corresponding to a given value to be found at a node (can aggregate locally if needed), but somewhat susceptible to skew
 - Range-based partitioning is very good to allow for all tuples corresponding to a value or range to be found at a node (can aggregate locally if needed), but very susceptible to skew
 - Round-robin is not susceptible to skew but has no locality benefits. All nodes need to participate in every query.



Replication

- In some cases, there may be benefits to not just partition the data but replicate it across nodes
 - Allowing for failures to some nodes to not affect the progress of queries
 - If one node is “busy” other node with same data can help with that query
- However: same issue as with materialization — need to keep replicas in sync to avoid staleness issues
- More on transactional issues possibly later (not a focus of this class)



Other Query Optimization Concerns

- Indexes:
 - Can be constructed globally (which partition contains the relevant data) and locally (which block within a partition contains the relevant data)
- Parallel versions of scans, sorts, joins, ...
 - Scans are easy — scan each partition in parallel
 - Q: how do we sort in parallel?



Parallel Sort

- Recall regular sort:
 - Pass 1:
 - Sort $B(R)/M$ subsets of M blocks of R each
 - Each is output to disk as a *run*
 - Pass 2:
 - “Merge” these runs by bringing in one block for each of them
- Instead, for pass 1: construct *partitioned runs* at each node
 - e.g., node 1 creates runs from [1-100], [101-200]...
- Pass 2: merge these partitioned runs across nodes
 - Node 1 gets all runs corresponding to [1-100] and then then merges them
 - Node 2 gets all runs corresponding to [101-200] and then then merges them
 - ...
- Pass 3: The sorted runs at Node 1, Node 2, ... are kept as is in a partitioned manner, or are concatenated to give an overall sort

